

# Introduction to Convolutional Neural Network

2019 - 2020 - 2021

Ando Ki, Ph.D.

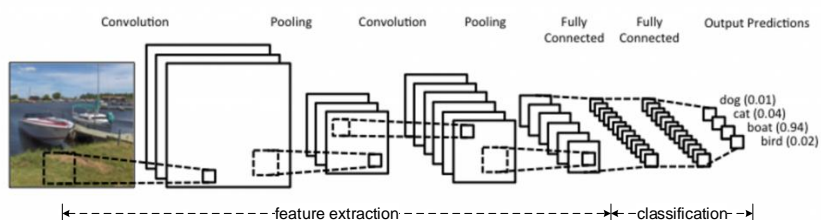
[adki@future-ds.com](mailto:adki@future-ds.com)

## Table of contents

- CNN
- CNN: convolution
- CNN: pooling
- CNN abstraction
- CNN examples

# CNN: Convolutional Neural Network

- CNN is a neural network that uses convolution in place of general matrix multiplication in at least one of their layers.
- General form of CNN (Convolutional Neural Network) for image classification
  - Feature extraction
    - ➔ Convolution
    - ➔ Pooling (sub-sampling)
  - Classification
    - ➔ Regression




Copyright (c) Ando Ki

CNN

3


## CNN: convolution

- Convolution is a matrix multiplication



When the value is large after convolution, it means there is a feature about it.

- It can be seen as a feature extractor



<https://adeshpande3.github.io/adeshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks/>


Copyright (c) Ando Ki

CNN

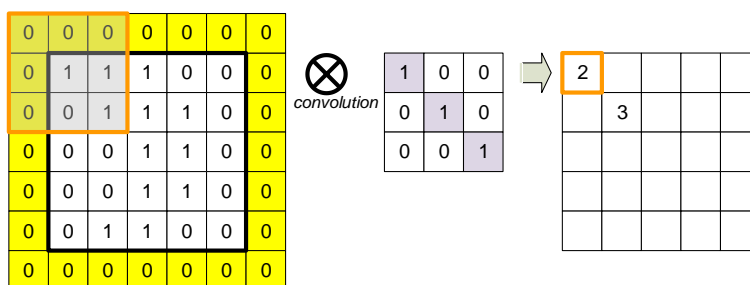
4

## CNN: convolution padding

- No padding
  - ▶ **Valid padding**



- Zero padding
  - ▶ **Same padding** due to input and out have the same dimensions.




Copyright (c) Ando Ki

CNN

5

## CNN: Pooling

- Pooling, i.e., sub-sampling
  - ▶ Max pooling
  - ▶ Average pooling



Copyright (c) Ando Ki

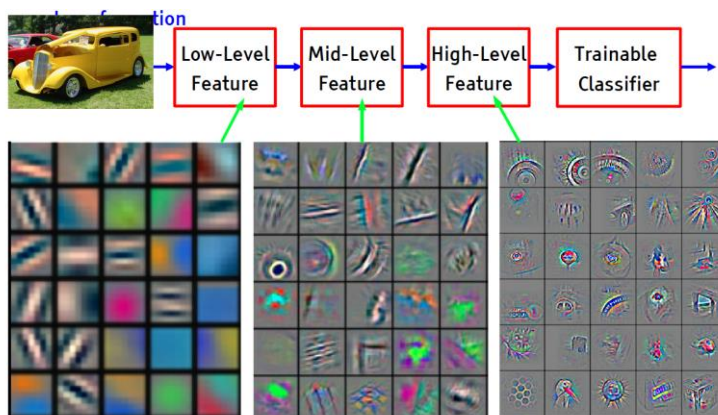
CNN

6

## How to choose filters

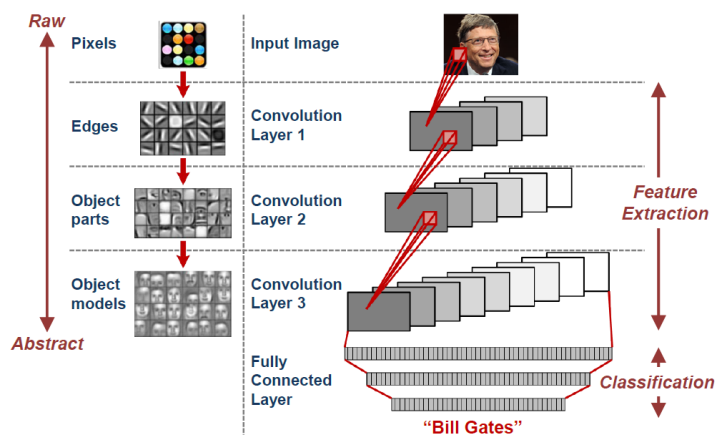
- With CNN/ConvNet the goal is to learn the filters; you don't actually design these filters (or kernels). They will be learned during training as long as the training converges.
- Initializing these filter parameters with good defaults before starting the training is key to convergence especially in very deep networks.
- Convolution filters can be initialized in one of the following ways.
  - ▶ 1. Randomly assigning weights for the different filters.
  - ▶ 2. Handcrafting the weights of the different filters to detect specific features during convolution.
  - ▶ 3. Learning filter weights using unsupervised training schemes.

## Deep learning: Learning Hierarchical Representations

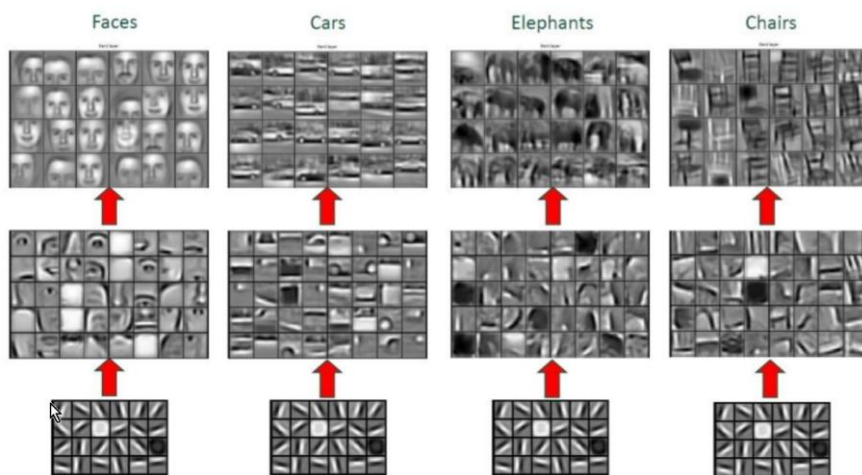


Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

# CNN abstraction

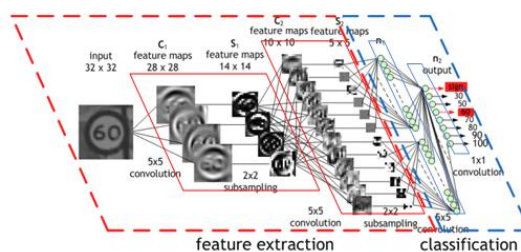
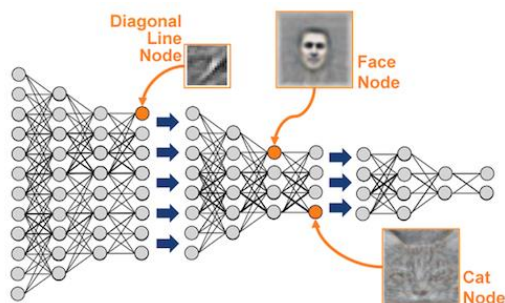


# CNN abstraction



Convolutional deep belief networks for scalable unsupervised learning of hierarchical representation". Lee et al., 2012

## CNN examples


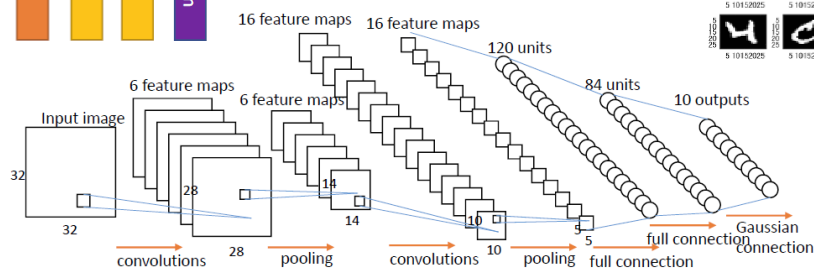



Copyright (c) Ando Ki

CNN

11

## CNN examples: LeNet-5 (1989)



Y. LeCun, B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel. Handwritten digit recognition with a back-propagation network. NIPS 1989.

Copyright (c) Ando Ki

CNN

12

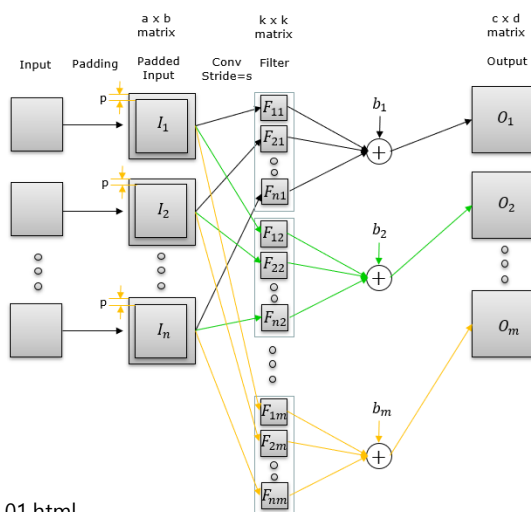
## Table of contents

- Convolution by matrix multiplication
- Fully connected layer by convolution
- Convolution for multi-channel
- Convolution and deconvolution
- Standard convolution
- Separable convolution: spatially separable
- Separable convolution: depthwise separable

## Convolution 2D

`torch.nn.Conv2d(in_channels=n,  
out_channels=m, kernel_size=k,  
stride=s, padding=p)`


*Note that there are  $n \times m$  filters (kernels)*




[http://sharetechnote.com/html/Python\\_PyTorch\\_nn\\_conv2D\\_01.html](http://sharetechnote.com/html/Python_PyTorch_nn_conv2D_01.html)

# Standard convolution

## Single channel with single kernel



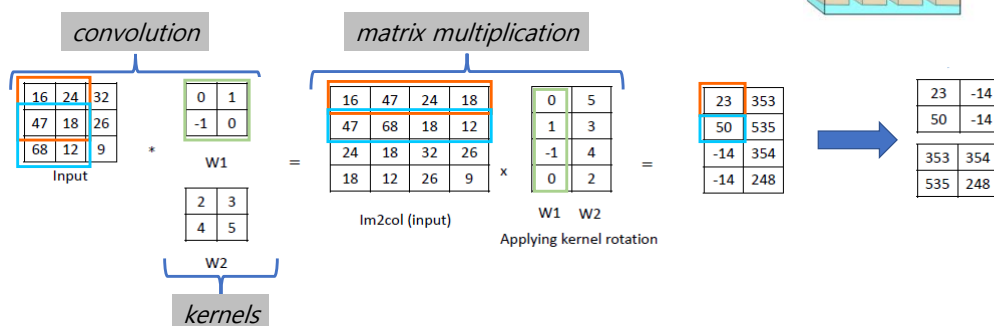
## Multiple channels and kernels



# Convolution by matrix multiplication


## 1-channel 2D convolution example: 3x3 input with 2x2 two kernel case

► Refer to 'im2col'





# im2col




Copyright (c) Ando Ki

CNN

17

## Fully connected layer by convolution (1/2)

### Convolution with kernels equal to the input size



Fully connected layer

remember, these also involve dot products between the receptive fields and kernels



where  $\mathbf{W}_1 = \begin{bmatrix} w_{1,1} & w_{1,2} \\ w_{1,3} & w_{1,4} \end{bmatrix}$

$\mathbf{W}_2 = \begin{bmatrix} w_{2,1} & w_{2,2} \\ w_{2,3} & w_{2,4} \end{bmatrix}$

```
conv = torch.nn.Conv2d(in_channels=1,
                        out_channels=2,
                        kernel_size=inputs.squeeze(dim=(0)).squeeze(dim=(0)).size())
print(conv.weight.size())
print(conv.bias.size())
```


Copyright (c) Ando Ki

CNN


18

## Fully connected layer by convolution (2/2)

### ■ Convolution with 1x1 kernels



Fully connected layer



Or, we can concatenate the inputs into 1x1 images with 4 channels and then use 2 kernels (remember, each kernel then also has 4 channels)


```
conv = torch.nn.Conv2d(in_channels=4,
                       out_channels=2,
                       kernel_size=(1, 1))
```

```
conv.weight.data = weights.view(2, 4, 1, 1)
conv.bias.data = bias
torch.relu(conv(inputs.view(1, 4, 1, 1)))
```

## Convolution for multi-channel


### ■ 2D convolution

- single and multi-channel




### ■ 3D convolution

- single and multi-channel



## Convolution and deconvolution

- Standard convolution (discrete convolution)
  - ▶ to extract feature map
- Standard deconvolution
  - ▶ known as **transposed convolution**
  - ▶ to reconstruct original image
  - ▶ a reverse operation of convolution



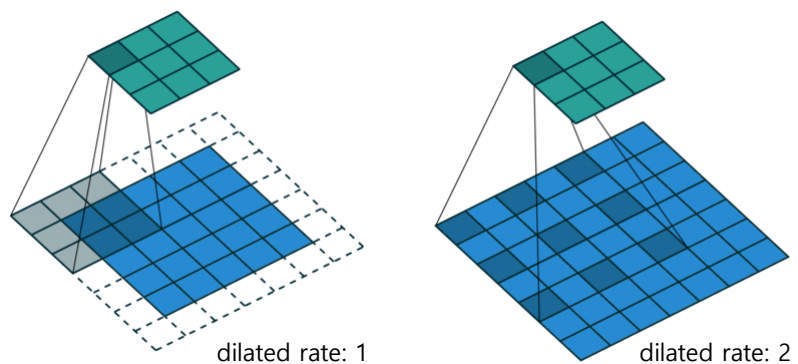
Copyright (c) Ando Ki

CNN

21

## Dilated convolution (atrous convolutions)

- Similar with deconvolution used in real-time segmentation
- smaller kernel for wider view
- not reverse operation (i.e, not reconstruction of original image)



Copyright (c) Ando Ki


CNN

22

## Separable convolution: spatially separable


### ■ standard convolution

- ▶ multiplications
  - $K \cdot K \cdot W \cdot W$



### ■ spatially separable convolution

- ▶ multiplications
  - $K \cdot W \cdot W + K \cdot W \cdot W$
  - 2/K ratio comparing to standard convolution



### ■ kernel divided


$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 3 & 6 & 9 \\ 4 & 8 & 12 \\ 5 & 10 & 15 \end{bmatrix} = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 3 \end{bmatrix}$$


Although spatially separable convolutions save cost, it is rarely used in deep learning. One of the main reason is that not all kernels can be divided into two, smaller kernels.

## Separable convolution: depthwise separable


### ■ Standard convolution and depthwise separable (no channel-wise conv)



Standard conv.

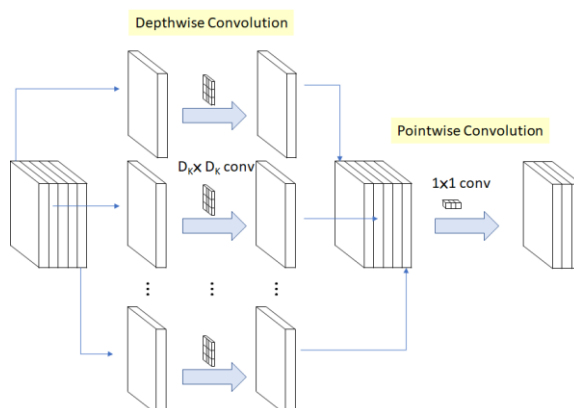


Depth-wise conv.



Point-wise conv.


### ■ MobileNet case



# Separable convolution: depthwise separable

## ■ Standard convolution

- uses kernels of a number of output channels



## ■ Depth wise separable

- Depthwise convolution: filtering stage
  - ⦿ uses kernels of a number of input channels
- Pointwise convolution: combining state
  - ⦿ uses kernels of a number of output channels

