

LeNet-5 Introduction

- a neural network architecture for handwritten and machine-printed character recognition -

2020 - 2021

Ando Ki, Ph.D.


adki@future-ds.com

Table of contents

- MNIST
- LeNet-5 for MNIST
- LeNet-5 layer 1-7
- References


MNIST

- MNIST (Modified National Institute of Standards and Technology database)
 - ▶ Modified National Institute of Standards and Technology
 - ▶ Handwritten digits database
 - ➔ 10 classes: 0, 1, ..., 9
 - ➔ training set: 60,000 training image
 - ➔ test set: 10,000 testing image




Note that MNIST images are 28x28.

Why background color is black?




LeNet-5 for MNIST

- LeNet is one of the popular convolutional networks, and works well on digit classification tasks.
 - ➔ 1024 (32x32) inputs of black and white → converted to floating number 0.0 ~ 1.0
 - ➔ 10 outputs representing digit 0 to 9



LeNet-5 for MNIST



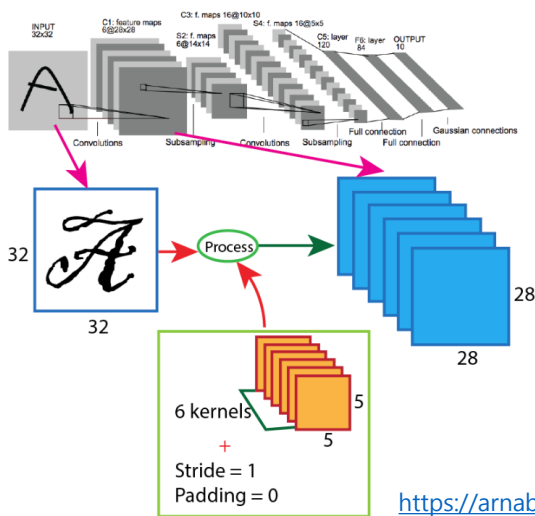
Summary of LeNet-5 Architecture

Original Image published in [LeCun et al., 1998]

	Layer	Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	32x32	-	-	-
1	Convolution	6	28x28	5x5	1	tanh
2	Average Pooling	6	14x14	2x2	2	tanh
3	Convolution	16	10x10	5x5	1	tanh
4	Average Pooling	16	5x5	2x2	2	tanh
5	Convolution	120	1x1	5x5	1	tanh
6	FC	-	84	-	-	tanh
Output	FC	-	10	-	-	softmax

tanh: hyperbolic tangent

LeNet-5 for MNIST: layer 1

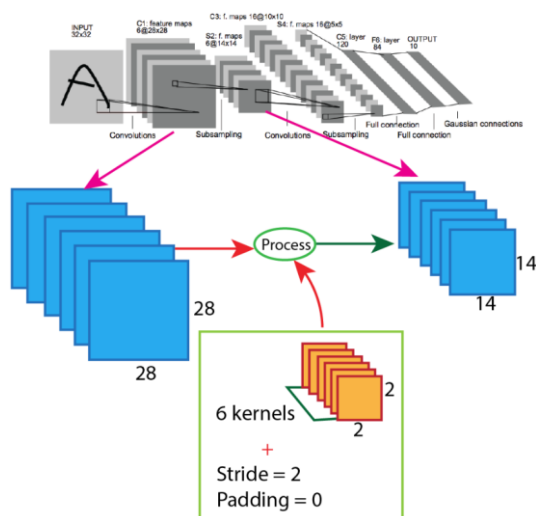


- 1st convolution layer
 - Input: 32x32 pixels (W=32)
 - Convolution filter: 6 kernels with 5x5 (K=5)
 - Parameters: weight+bias=1x5x5x6+6=156
 - Convolution padding: 0 (P=0)
 - Convolution: stride 1 (S=1)
 - Activation: ReLU (Tanh)
 - Results in: 6 features of 28x28

$$Output = \frac{W - K + 2(P)}{S} + 1 = \frac{32 - 5 + 2(0)}{1} + 1 = 28$$

https://arnabfly.github.io/arnab_blog/lenet5/

LeNet-5 for MNIST: layer 2




- 1st pooling layer (sub-sampling)
 - Input: 6 features with 28x28 (W=28)
 - Max pooling filter: 2x2 (K=2)
 - Pooling padding: 0 (P=0)
 - Pooling: stride 2 (S=2)
 - It generates $\frac{1}{2}$ number of elements
 - Activation: ReLU
 - Results in: 6 features of 14x14

$$\text{Output} = \frac{W - K + 2(P)}{S} + 1 = \frac{28 - 2 + 2(0)}{2} + 1 = 14$$

Copyright (c) Ando Ki

7

LeNet-5 for MNIST: layer 3



- 2nd convolution
 - Input: 6 features with 14x14 pixels (W=14)
 - 6 kernels are used at the previous stage
 - Convolution filter: 16 kernels with 5x5 (K=5)
 - Parameters: $6 \times 5 \times 5 \times 16 = 2,416$
 - Reduction: $6 \times 5 \times 5 \times 10 = 1,516$
 - Convolution padding: 0 (P=0)
 - Convolution: stride 1 (S=1)
 - It generates the same number of elements
 - Activation: ReLU
 - Results in: 16 features of 10x10


		feature maps															
input		0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	0	X				X	X	X			X	X	X	X		X	X
	1	X	X				X	X	X			X	X	X	X		X
	2	X	X	X				X	X	X			X		X	X	X
	3		X	X	X			X	X	X	X			X		X	X
	4			X	X	X			X	X	X	X		X	X	X	X

only 10 out of 16 feature maps are connected to each kernel.

Copyright (c) Ando Ki

8

LeNet-5 for MNIST: layer 4




- 2nd pooling
 - ▶ Input: 16 features with 10x10
 - ▶ Max pooling filter: 2x2
 - ▶ Pooling padding: 0
 - ▶ Pooling: stride 2
 - ⇒ It generates $\frac{1}{2}$ number of elements
 - ▶ Results in: 16 features of 5x5

Copyright (c) Ando Ki

9

LeNet-5 for MNIST: layer 5




- fully connected layer for flatten
 - ▶ Input: 16 features with 5x5
 - ▶ Reshaping: 3-D array to 1-D vector
 - ⇒ $16 \times 5 \times 5 \rightarrow 400$
 - ▶ Output: 120
 - ⇒ Neurons: 120
 - ▶ Parameters (weights+bias)
 - ⇒ $400 \times 120 + 120 = 48,120$

Copyright (c) Ando Ki

10

LeNet-5 for MNIST: layer 6

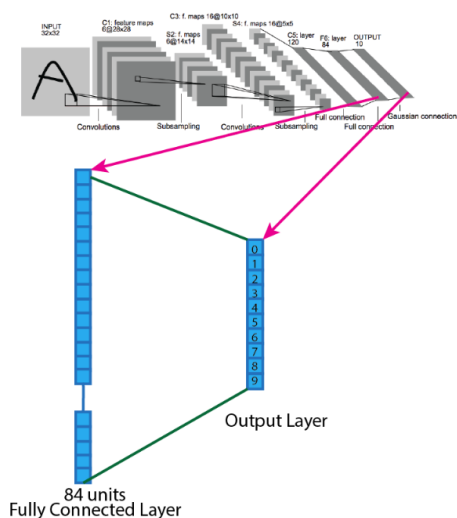


- fully connected layer for flatten
 - ▶ Input: 120 feature map
 - ▶ Output: 84
 - ⇒ Neurons: 84
 - ▶ Parameters: $120 \times 84 + 84 = 10,164$

Copyright (c) Ando Ki

11

LeNet-5 for MNIST: layer 7



- fully connected layer for flatten
 - ▶ Input: 84 feature map
 - ▶ Output: 10
 - ⇒ Neurons: 10
 - ▶ Parameters: $84 \times 10 + 10 = 850$

The output layer is composed of Euclidean Radial Basis Function unit (RBF), one for each class, with 84 inputs each. The outputs of each RBF unit i -th y is computed as follows.

$$y_i = \sum_j (x_j - w_{ij})^2.$$

Nowadays, softmax is used instead.

Copyright (c) Ando Ki


12

Summary of LeNet-5


Summary of LeNet-5 Architecture

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	32x32	-	-	-
1	Convolution	6	28x28	5x5	1	tanh
2	Average Pooling	6	14x14	2x2	2	tanh
3	Convolution	16	10x10	5x5	1	tanh
4	Average Pooling	16	5x5	2x2	2	tanh
5	Convolution	120	1x1	5x5	1	tanh
6	FC	-	84	-	-	tanh
Output	FC	-	10	-	-	softmax

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 6)	156
average_pooling2d (AveragePo	(None, 14, 14, 6)	0
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2416 (1516)
average_pooling2d_1 (Average	(None, 5, 5, 16)	0
flatten (Flatten)	(None, 400)	0
dense (Dense)	(None, 120)	48120
dense_1 (Dense)	(None, 84)	10164
dense_2 (Dense)	(None, 10)	850
Total params: 61,706		
Trainable params: 61,706		
Non-trainable params: 0		




Running LeNet-5: Darknet case



LeNet-5 configuration

[net] batch=100 subdivisions=1 height=32 width=32 channels=1 momentum=0.9 decay=0.00005 max_crop=28 learning_rate=0.01 policy=poly power=4 max_batches=500	angle=1 hue=1 saturation=1 exposure=1 aspect=1 [convolutional] filters=6 size=5 stride=1 pad=0 activation=relu	[maxpool] size=2 stride=2 [convolutional] filters=16 size=5 stride=1 pad=0 activation=relu [maxpool] size=2 stride=2	[connected] output= 120 activation=relu [connected] output= 84 activation=relu [connected] output= 10 activation=linear	[softmax] groups=1 [cost] type=sse
--	--	---	---	---

Running LeNet-5: PyTorch case (1/2)




```
from torch import nn
from torch.nn import Module
import torch.nn.functional as F

class Lenet5Model(Module):
    def __init__(self):
        super(Lenet5Model, self).__init__()
        self.conv1 = nn.Conv2d( in_channels=1, out_channels=6, kernel_size=(5,5)
                                , stride=1, padding=0, bias=True)
        self.relu1 = nn.ReLU()
        self.pool1 = nn.MaxPool2d( kernel_size=(2,2), stride=2 )
        self.conv2 = nn.Conv2d( in_channels=6, out_channels=16
                                , kernel_size=(5,5), stride=1
                                , padding=0
                                , bias=True)
        self.relu2 = nn.ReLU()
        self.pool2 = nn.MaxPool2d( kernel_size=(2,2), stride=2 )
```

Copyright (c) Ando Ki

15

Running LeNet-5: PyTorch case (2/2)



```
self.fc1 = nn.Linear( in_features=16*5*5
                      , out_features=120
                      , bias=True)
self.relu3 = nn.ReLU()
self.fc2 = nn.Linear( in_features=120
                      , out_features=84
                      , bias=True)
self.relu4 = nn.ReLU()
self.fc3 = nn.Linear( in_features=84
                      , out_features=10
                      , bias=True)
self.relu5 = nn.ReLU()
```


Copyright (c) Ando Ki

16

LeNet-5 for MNIST: layer


1st convolution layer

- ▶ Input: 28x28 pixels
- ▶ Convolution filter: 32 kernels with 5x5
- ▶ Convolution: stride 1
 - ☞ It generates the same number of elements
- ▶ Results in: 32 features of 28x28



1st pooling layer (sub-sampling)

- ▶ Input: 32 features with 28x28
- ▶ Max pooling filter: 5x5 (2x2 ?)
- ▶ Convolution: stride 2
 - ☞ It generates $\frac{1}{2}$ number of elements
- ▶ Results in: 32 features of 14x14




Copyright (c) Ando Ki

17

LeNet-5 for MNIST: layer


2nd convolution

- ▶ Input: 32 features with 14x14 pixels
 - ☞ 32 kernels are used at the previous stage
- ▶ Convolution filter: 64 kernels with 5x5
- ▶ Convolution: stride 1
 - ☞ It generates the same number of elements
- ▶ Results in: 64 features of 14x14



2nd pooling

- ▶ Input: 64 features with 14x14
- ▶ Max pooling filter: 2x2
- ▶ Convolution: stride 2
 - ☞ It generates $\frac{1}{2}$ number of elements
- ▶ Results in: 64 features of 7x7



Copyright (c) Ando Ki

18

LeNet-5 for MNIST: layer

fully connected layer

- ▶ Input: 64 features with 7x7
- ▶ Reshaping: 3-D array to 1-D vector
 - ◉ $64 \times 7 \times 7 \rightarrow 3,136$
- ▶ Neurons: 1024

read-out layer

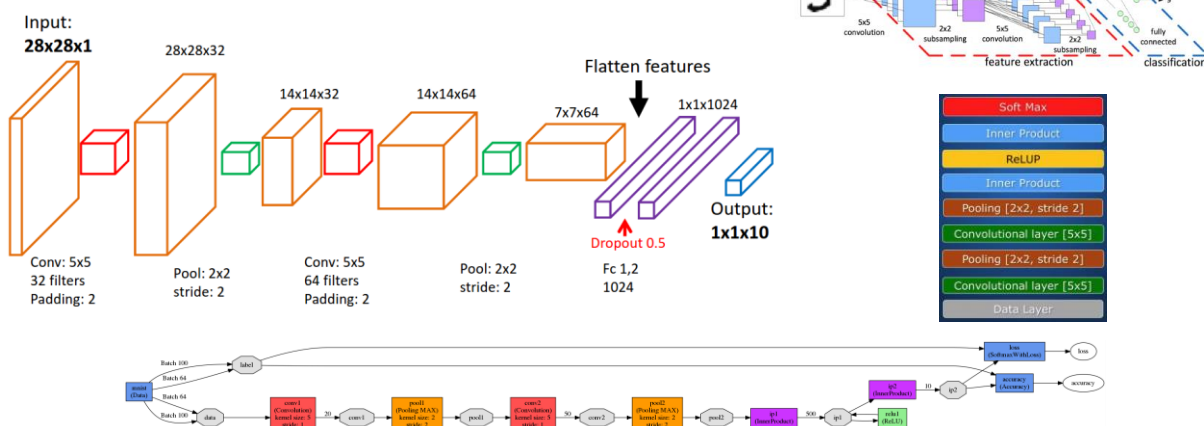
- ▶ Input: 1024 neurons
- ▶ Output: 10 classes



Copyright (c) Ando Ki

19

LeNet-5 for MNIST: all together



Copyright (c) Ando Ki

20

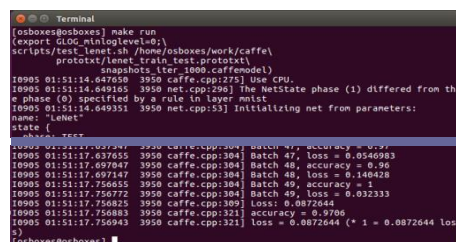
LeNet-5 for MNIST: running

Steps (in details)

- ▶ go to project directory
 - `$ cd work/codes/caffe_v1-projects/mnist.LeNet`
- ▶ get dataset:
 - `$./scripts/get_mnist.sh data`
 - (ungzip all in 'data' directory)
- ▶ convert the dataset to Caffe data format
 - `$./scripts/create_mnist.sh ${CAFFE_HOME} data`
- ▶ training
 - `$./scripts/train_lenet.sh ${CAFFE_HOME} prototxt/lenet_solver.prototxt`
- ▶ running LeNet model with 'mnist_test_lmdb'
 - `$./scripts/test_lenet.sh`

Step in simple

- ▶ go to project directory
 - `$ cd work/codes/caffe_v1-projects/mnist.LeNet`
- ▶ Run make
 - `$ make cleanall`
 - `$ make lmdb`
 - `$ make train`
 - `$ make test`



```

[osboxes@osboxes] make run
[export GLIBC_TUNABLES=glibc.libc.so.6]
[scripts/test_lenet.sh /home/osboxes/work/caffe/prototxt/lenet_train_test.prototxt]
[0905 01:51:14.047650 3950 caffe.cpp:275] Use CPU.
[0905 01:51:14.049105 3950 net.cpp:290] The NetState phase (1) differed from the phase (0) specified by a rule in layer mnist
[0905 01:51:14.049351 3950 net.cpp:53] Initializing net from parameters:
name: "LeNet"
state {
  phase: "train"
}
[0905 01:51:17.037341 3950 caffe.cpp:304] Batch 47, accuracy = 0.957
[0905 01:51:17.637055 3950 caffe.cpp:304] Batch 47, loss = 0.0546983
[0905 01:51:17.697047 3950 caffe.cpp:304] Batch 48, accuracy = 0.956
[0905 01:51:17.697147 3950 caffe.cpp:304] Batch 48, loss = 0.140428
[0905 01:51:17.756855 3950 caffe.cpp:304] Batch 49, accuracy = 1
[0905 01:51:17.756772 3950 caffe.cpp:304] Batch 49, loss = 0.032333
[0905 01:51:17.756825 3950 caffe.cpp:309] Loss: 0.0872644
[0905 01:51:17.756983 3950 caffe.cpp:321] accuracy = 0.9706
[0905 01:51:17.756943 3950 caffe.cpp:321] loss = 0.0872644 (* 1 = 0.0872644 loss)
[osboxes@osboxes]

```

LeNet-5 for MNIST: solver

- ▶ net: network mode
- ▶ test_iter: iterations to test
- ▶ test_interval: interval between test
- ▶ base_lr: Learning Rate initial value
- ▶ display: iterations to show progress
- ▶ max_iter: max iterations for training.
- ▶ snapshot: iterations to store snapshot.
- ▶ solver_mode: CPU or GPU

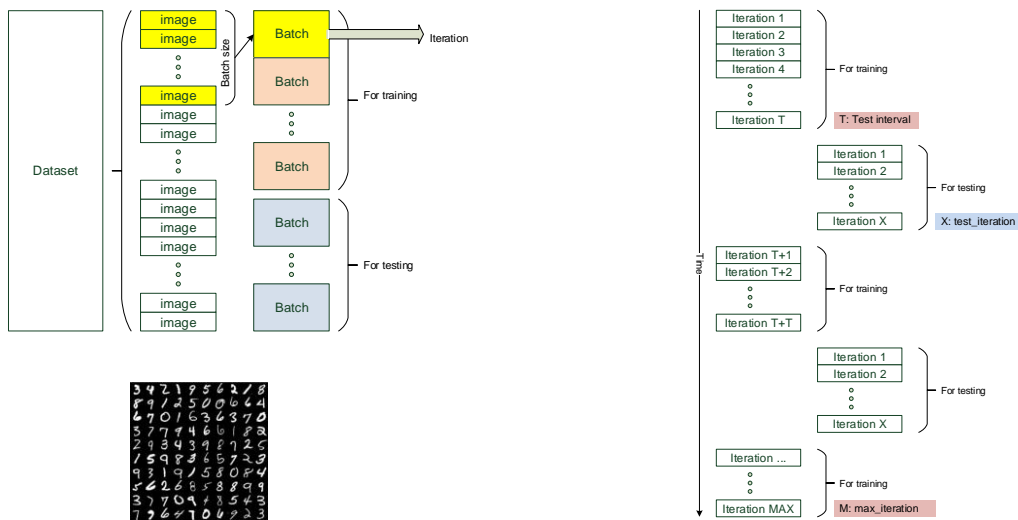
```

# MNIST lenet_solver.prototxt
net: "lenet_train_test.prototxt"
test_iter: 100
test_interval: 500
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
lr_policy: "inv"
gamma: 0.0001
power: 0.75
display: 100
max_iter: 10000
snapshot: 5000
snapshot_prefix: "snapshots"
solver_mode: CPU

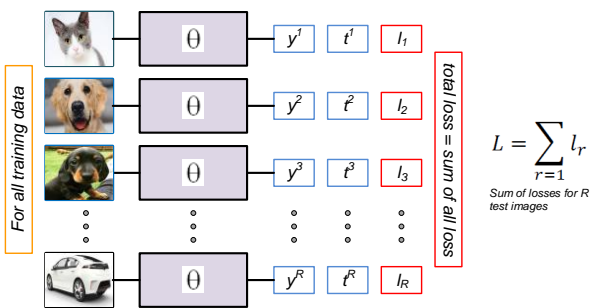


```

<https://github.com/BVLC/caffe/wiki/Solver-Prototxt>

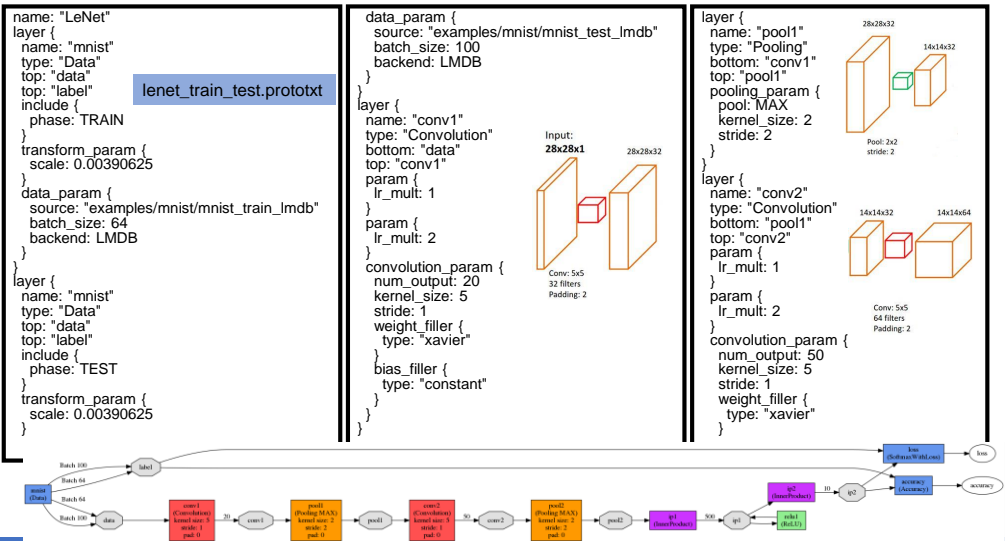
LeNet-5 for MNIST: solver



LeNet-5 for MNIST: solver



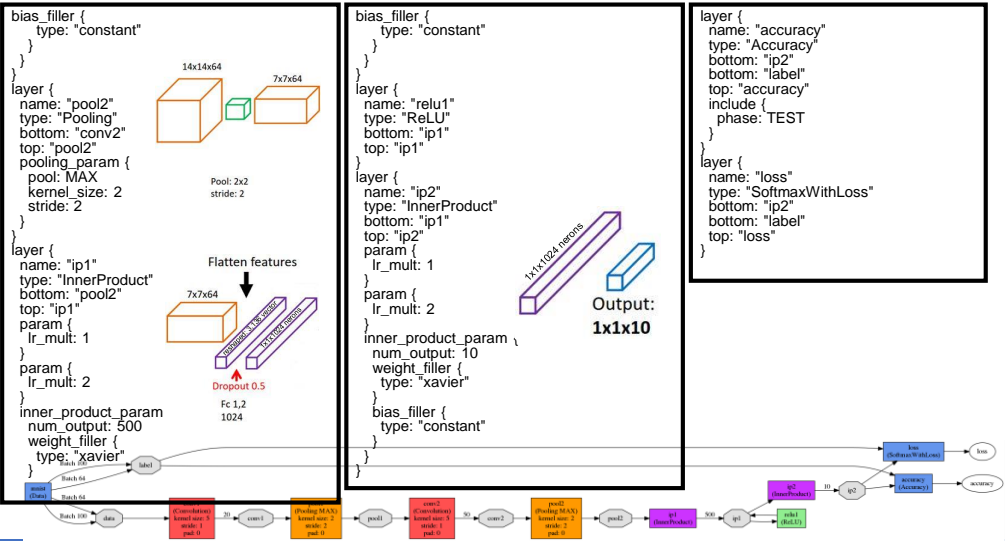
LeNet-5 for MNIST: net



Copyright (c) Ando Ki

25

LeNet-5 for MNIST: net




Copyright (c) Ando Ki

26

Running LeNet with Caffe

- This example is about LeNet
- Make sure 'work/caffe' is ready
 - ▶ see the pervious slides
 - ▶ Step 1: go to your project directory
 - ⇒ [user@host] cd \$(PROJECT)/codes.caffe/mnist.LeNet
 - ▶ Step 2: check network
 - ⇒ [user@host] make draw
 - ⇒ [user@host] fim lenet_train_test.png
 - ▶ Step 3: make data (convert data)
 - ⇒ [user@host] make lmdb
 - ▶ Step 4: run train (it takes time)
 - ⇒ [user@host] make train
 - ▶ Step 5: run loss graph
 - ⇒ [user@host] make plot
 - ▶ Step 6: run test
 - ⇒ [user@host] make test
 - ▶ Step 7: run deployment (inference)
 - ⇒ [user@host] make deploy



```

[osboxes@osboxes] make run
(export: LOG_minloglevel=0)
scripts/test_lenet.sh /home/osboxes/work/caffe/
prototxt/lenet_train_test.prototxt\
snapshots_iter_1000.caffemodel)
10905 01:51:14.647650 3950 caffe.cpp:275] Use CPU.
10905 01:51:14.649165 3950 net.cpp:296] The NetState phase (1) differed from th
e phase (0) specified by a rule in layer mnist
10905 01:51:14.649351 3950 net.cpp:53] Initializing net from parameters:
name: "LeNet"
state {
  ...
  TEST
10905 01:51:17.037341 3950 caffe.cpp:304] Batch 47, accuracy = 0.97
10905 01:51:17.637655 3950 caffe.cpp:304] Batch 47, loss = 0.0540983
10905 01:51:17.697047 3950 caffe.cpp:304] Batch 48, accuracy = 0.96
10905 01:51:17.697147 3950 caffe.cpp:304] Batch 48, loss = 0.140428
10905 01:51:17.756655 3950 caffe.cpp:304] Batch 49, accuracy = 1
10905 01:51:17.756772 3950 caffe.cpp:304] Batch 49, loss = 0.032333
10905 01:51:17.756825 3950 caffe.cpp:309] Loss: 0.0872644
10905 01:51:17.756833 3950 caffe.cpp:321] accuracy = 0.9766
10905 01:51:17.756943 3950 caffe.cpp:321] loss = 0.0872644 (* 1 = 0.0872644 los
s)
[osboxes@osboxes]

```

use 'display' for Ubuntu, 'fim' for Raspbian to display image.


Run inference with sample image

- Go to 'mnist.LeNet' directory
 - ▶ \$ cd .../codes/caffe_v1-project/mnist.LeNet
- Run make
 - ▶ \$ make deploy
- Note that LeNet uses inverted image, i.e., background should be black.

```

[osboxes@osboxes] make
LOG_minloglevel=2 python mnist_test.py samples/4.png
[ 2.50871176e-06 1.81367841e-06 1.57160230e-05 2.03305008e-05
 9.89796937e-01 9.22584604e-06 5.09644167e-07 4.18145180e-04
 5.32380818e-06 9.72963311e-03]
4
[osboxes@osboxes]

```



Deploy prototxt (1/2)

■ Refer to 'lenet_deploy.prototxt' under 'prototxt' directory.

```
##### Remove the data layer
#layer {
#  name: "mnist"
#  type: "Data"
#  top: "data"
#  top: "label"
#  include {
#    phase: TRAIN
#  }
#  transform_param {
#    scale: 0.00390625
#  }
#  data_param {
#    source: "data/mnist_train_lmdb"
#    batch_size: 64
#    backend: LMDB
#  }
#}

##### Remove the label layer
#layer {
#  name: "mnist"
#  type: "Data"
#  top: "data"
#  top: "label"
#  include {
#    phase: TEST
#  }
#  transform_param {
#    scale: 0.00390625
#  }
#  data_param {
#    source: "data/mnist_test_lmdb"
#    batch_size: 100
#    backend: LMDB
#  }
#}
```

Remove

Remove

```
##### Add a new layer to accept data without label
### It define the name and shapes of the input blobs.
#  shape: { dim: 1 # batchsize (how many images/samples are fed through the
#               # network in parallel)
#           dim: 1 # number of channels; 1 means greay, not RGB
#           dim: 28 # height of data, i.e., pixels (MNIST data is 28x28)
#           dim: 28 # width of data, i.e., pixels (MNIST data is 28x28)
#         }
#  layer {
#    name: "data"
#    type: "input"
#    top: "data"
#    input_param {
#      shape: { dim: 1
#               dim: 1
#               dim: 28
#               dim: 28
#             }
#    }
#  }
#}

layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
}

#####
other hidden layers remains
```

Add

Deploy prototxt (2/2)

■ Refer to 'lenet_deploy.prototxt' under 'prototxt' directory.

```
...
...
}
}

##### Remove the layers depending upon data labels
##### accordingly remove layer that uses 'data' as bottom
#layer {
#  name: "accuracy"
#  type: "Accuracy"
#  bottom: "ip2"
#  bottom: "label"
#  top: "accuracy"
#  include {
#    phase: TEST
#  }
#}

#layer {
#  name: "loss"
#  type: "SoftmaxWithLoss"
#  bottom: "ip2"
#  bottom: "label"
#  top: "loss"
#}

##### Add a new layer to the end of this network to produce a Softmax
output
layer {
  name: "loss"
  type: "Softmax"
  bottom: "ip2"
  top: "loss"
}
```

Remove

Remove

Add

Run inference with sample image (another way)


- Go to 'mnist.LeNet.python' directory

- ▶ \$ cd ../codes/caffe_v1-project/mnist.LeNet.python

- Run make

- ▶ \$ make

```
Terminal
[osboxes@osboxes] make
GLOG_minloglevel=2 python mnist_test.py samples/4.png
[ 2.50871176e-06  1.81367841e-06  1.57160230e-05  2.03305008e-05
 9.89796937e-01  9.22584604e-06  5.09644167e-07  4.18145180e-04
 5.32380818e-06  9.72963311e-03]
4
[osboxes@osboxes]
```



Caffe Python interface for LeNet

```
import os
os.environ['GLOG_minloglevel']='2'

import caffe
import numpy as np
import cv2
import sys
import Image

caffe_home = os.environ['CAFFE_ROOT'];
model = caffe_home + '/examples/mnist/lenet.prototxt';
weights = './mnist.LeNet/snapshots_iter_1000.caffemodel';
net = caffe.Net(model, weights, caffe.TEST);
caffe.set_mode_cpu()

img = cv2.imread(sys.argv[1],0)
if img.shape != [28,28]:
    img2 = cv2.resize(img,(28,28))
    img = img2.reshape(28,28,-1);
else:
    img = img.reshape(28,28,-1);

img = 1.0 - img/255.0

out = net.forward_all(data=np.asarray([img.transpose(2,0,1)]))

print out['prob'][0]
print out['prob'][0].argmax()
```

Using pycaffe

This may not need.

Revert image and normalize it to 0~1

Inference

Get the highest probability one

Running LeNet-5: TensorFlow case

Running LeNet-5: Keras case

References

- Yann LeCun and et.al., Gradient-Based Learning Applied to Document Recognition, Proc. of the IEEE, Nov. 1998.
- Break Down Lenet-5
 - ▶ https://arnabfly.github.io/arnab_blog/lenet5/
- LeNet-5 – A Classic CNN Architecture
 - ▶ <https://www.datasciencecentral.com/profiles/blogs/lenet-5-a-classic-cnn-architecture>
 - ▶ <https://engmrk.com/lenet-5-a-classic-cnn-architecture/>