

Introduction to PyTorch

2020 - 2021

Ando Ki, Ph.D.

adki@future-ds.com

Table of contents

- What is PyTorch
- PyTorch components
- PyTorch is imperative program
- Three levels of abstraction
- PyTorch tensor v.s. NumPy ndarray
- PyTorch functions, dimensionality
- Building network
- An example of model
- Defining an optimizer
- Save/load models

What is PyTorch


- **PyTorch** is a machine learning Python library, developed by the Facebook AI research group.

- ▶ Based on Torch library

- ➔ Torch is an open-source machine learning library, a scientific computing framework, and a script language based on the Lua programming language.[3] It provides a wide range of algorithms for deep learning, and uses the scripting language LuaJIT, and an underlying C implementation.

- **PyTorch** is a Python-based scientific computing package targeted at two sets of audiences:

- ▶ A replacement for NumPy to use the power of GPUs → tensor
- ▶ a deep learning research platform that provides maximum flexibility and speed



Installing PyTorch

- Visit: <https://pytorch.org/get-started/locally/>
- Select your preferences and run the install command.
- Run the command
 - ▶ note 'conda' is required.

| | | | | |
|-------------------|---|------|-------------------|--------|
| PyTorch Build | Stable (1.6.0) | | Preview (Nightly) | |
| Your OS | Linux | Mac | Windows | |
| Package | Conda | Pip | LibTorch | Source |
| Language | Python | | C++ / Java | |
| CUDA | 9.2 | 10.1 | 10.2 | None |
| Run this Command: | <pre>conda install pytorch torchvision cpuonly -c pytorch</pre> | | | |

PyTorch components


| Component | Description | Package | Description |
|------------------------------------|---|----------------------------------|--|
| <code>torch</code> | a Tensor library like NumPy, with strong GPU support | <code>torch</code> | The top-level PyTorch package and tensor library. |
| <code>torch.autograd</code> | a tape-based automatic differentiation library that supports all differentiable Tensor operations in torch | <code>torch.nn</code> | A subpackage that contains modules and extensible classes for building neural networks. |
| <code>torch.jit</code> | a compilation stack (TorchScript) to create serializable and optimizable models from PyTorch code | <code>torch.autograd</code> | A subpackage that supports all the differentiable Tensor operations in PyTorch. |
| <code>torch.nn</code> | a neural networks library deeply integrated with autograd designed for maximum flexibility | <code>torch.nn.functional</code> | A functional interface that contains typical operations used for building neural networks like loss functions, activation functions, and convolution operations. |
| <code>torch.multiprocessing</code> | Python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and Hogwild training | <code>torch.optim</code> | A subpackage that contains standard optimization operations like SGD and Adam. |
| <code>torch.utils</code> | DataLoader and other utility functions for convenience | <code>torch.utils</code> | A subpackage that contains utility classes like data sets and data loaders that make data preprocessing easier. |
| | | <code>torchvision</code> | A package that provides access to popular datasets, model architectures, and image transformations for computer vision. |

PyTorch is imperative program

- Imperative program
 - ▶ E.g., NumPy, PyTorch
 - ▶ When the line 'c=b*a' is executed, it actually runs a computation.

```
import numpy as np
a = np.ones(10)
b = np.ones(10) * 2
c = b * a
d = c + 1
```


- Symbolic program
 - ▶ E.g, MXNET, TensorFlow
 - ▶ When the line 'C=B*A' is executed, no computation occurs.
 - ▶ Instead, it generates a computation graph (i.e., a symbolic graph) that represents the computation.
 - ▶ It should be compiled and then run it.



Levels of abstraction

■ Tensor

- ▶ imperative ndarray can run on GPU/CPU (just like numpy array, but can run on GPU)
 - ➔ `x = torch.tensor([[1,2,3],[4,5,6]])`
- ▶ Variable (it is merged to 'Tensor')
 - ➔ a wrapper around PyTorch tensor, and represents a computational graph and stores data and gradient
 - ➔ class: "torch.autograd.Variable"
 - `x = torch.autograd.Variable(x, required_grad=False)`



■ Module

- ▶ a neural network layer and can store state or learnable weights

PyTorch tensor v.s. NumPy ndarray

- PyTorch's **tensors** are very similar to NumPy's **ndarrays**, but they have a device, 'cpu', 'cuda',

```
>>> t = torch.tensor([1,2,3], device='cpu',
... requires_grad=False, dtype=torch.float32)
>>> print(t.dtype)
torch.float32
>>> print(t.device)
cpu
```



```
>>> a = numpy.ones((3,4))
>>> t = torch.from_numpy(a) # get tensor

>>> b = torch.tensor([3,4])
>>> n = b.numpy() # get numpy

>>> c = torch.ones((3,4))
```



Torch data types

| Data Type | T.dtype | T.dtype alias | T.tensor() CPU alias | T.tensor() GPU alias |
|------------------------|-----------|---------------|----------------------|-----------------------|
| 16-bit floating point | T.float16 | T.half | T.HalfTensor() | T.cuda.HalfTensor() |
| 32-bit floating point | T.float32 | T.float | T.FloatTensor() | T.cuda.FloatTensor() |
| 64-bit floating point | T.float64 | T.double | T.DoubleTensor() | T.cuda.DoubleTensor() |
| 8-bit unsigned integer | T.uint8 | | T.ByteTensor() | T.cuda.ByteTensor() |
| 8-bit signed integer | T.int8 | | T.CharTensor() | T.cuda.CharTensor() |
| 16-bit signed integer | T.int16 | T.short | T.ShortTensor() | T.cuda.ShortTensor() |
| 32-bit signed integer | T.int32 | T.int | T.IntTensor() | T.cuda.IntTensor() |
| 64-bit signed integer | T.int64 | T.long | T.LongTensor() | T.cuda.LongTensor() |
| Boolean | T.bool | | T.BoolTensor() | T.cuda.BoolTensor() |

Default data type: torch.FloatTensor

Copyright (c) Ando Ki

9

PyTorch functions, dimensionality

- `x.size()` #* return tuple-like object of dimensions, old codes
- `x.shape` # return tuple-like object of dimensions, numpy style
- `x.ndim` # number of dimensions, also known as `.dim()`
- `x.view(a,b,...)` #* reshapes `x` into size `(a,b,...)`
- `x.view(-1,a)` #* reshapes `x` into size `(b,a)` for some `b`
- `x.reshape(a,b,...)` # equivalent with `.view()`
- `x.transpose(a,b)` # swaps dimensions `a` and `b`; only two dimensions
- `x.permute(*dims)` # permutes dimensions; missing in numpy; more than two dim.
- `x.unsqueeze(dim)` # tensor with added axis; missing in numpy
- `x.unsqueeze(dim=2)` # `(a,b,c)` tensor -> `(a,b,1,c)` tensor; missing in numpy

Copyright (c) Ando Ki

10

Tips for Torch

```
T = torch.tensor( data=x, dtype=torch.int32, device='cuda', required_grad='false')
```

- data: actual data in list or array
- dtype: data type
- device: cpu or cuda device
- required_grad: store gradient or not

how to create tensor

```
Y = torch.as_tensor(T, dtype=torch.half, device='cpu')  
Z = T.to(device='cpu', dtype=torch.float64)
```

how to change attributes

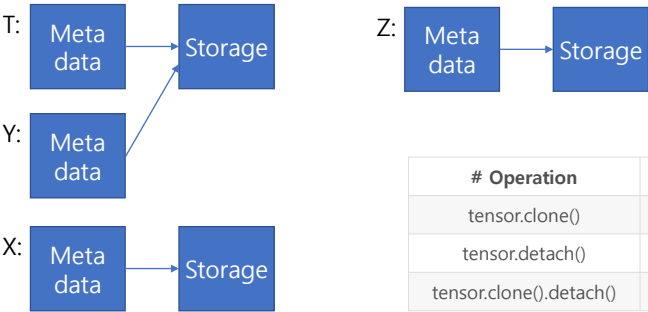
```
A = Z.abs()  
B = Z.abs_()
```

A and Z are different
B and Z shared the same data storage; make in-place tensor; mutable
(Note '_' underscore)

Tips for Torch

```
T = torch.tensor([1.,2.,3.], required_grad=True)  
X = T.clone()  
Y = T.detach()  
Z = T.clone().detach()
```

- clone() returns an identical tensor with new memory.
 - immutable
- detach() return an identical tensor with the same memory.
 - mutable
 - not involved gradient calculation




| # Operation | New/Shared memory | Still in computation graph |
|-------------------------|-------------------|----------------------------|
| tensor.clone() | New | Yes |
| tensor.detach() | Shared | No |
| tensor.clone().detach() | New | No |

Tips for Torch

```
>>> T = torch.zeros(3,2) # all zeroes
>>> X = T.view(2,3) # transpose style
>>> T.fill_(1) # in-place
tensor([[1,1],[1,1],[1,1]])
>>> X
tensor([[1,1,1],[1,1,1]])
```

- `view()` returns a new shape of tensor with the same memory.
 - mutable
 - working for contiguous storage
- `reshape()` returns a copy or a view of the original tensor.
 - working for contiguous or non-contiguous tensor.
 - use `clone()` if a copy is required.
 - use `view()` if the same storage.



Building network (1/3): using tensor

■ Two-layer net using PyTorch tensors.

- ▶ create random tensors for data and weights

```
import torch

dtype = torch.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)
```

- ▶ Forward pass
 - ➡ compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

- ▶ Backward pass
 - ➡ manually compute gradients

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)
```

- ▶ Gradient descent step on weights

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```

Building network (2/3): using Variable

Two-layer net using PyTorch Variable

- ▶ create random variables for data and weights

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10

x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)
```

- ▶ Forward pass
 - ➔ compute predictions and loss

```
learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

ReLU

SSE

- ▶ Backward pass
 - ➔ manually compute gradients

```
if w1.grad: w1.grad.data.zero_()
if w2.grad: w2.grad.data.zero_()
loss.backward()
```

- ▶ Gradient descent step on weights

```
w1.data -= learning_rate * w1.grad.data
w2.data -= learning_rate * w2.grad.data
```

Copyright (c) Ando Ki

15

Building network (3/3): using Variable

Building network using Variable with user defined Autograd function

- ▶ Define own autograd function by writing forward and backward for Tensors.

```
class ReLU(torch.autograd.Function):
    def forward(self, x):
        self.save_for_backward(x)
        return x.clamp(min=0)

    def backward(self, grad_y):
        x, = self.saved_tensors
        grad_input = grad_y.clone()
        grad_input[x < 0] = 0
        return grad_input
```

```
N, D_in, H, D_out = 64, 1000, 100, 10

x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    relu = ReLU()
    y_pred = relu(x.mm(w1)).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

Copyright (c) Ando Ki

16

Building network: creating the model using nn (1/2)

- A **model** is of a *nn.Module class* type. A model can contain other models.
 - ▶ The nn.Module's weights are called "Parameters".
 - ▶ A nn.Module consists of an **initialization** of the Parameters and a **forward function**.

```
class Model(nn.Module):
    def __init__(self):
        super().__init__()
        # structure definition and initialization
    def forward(self, x):
        # actual forward propagation
        result = processing(x)
        return result
```


You don't need to define a backward function!
 - Autograded variable will take care of it.

Forward is called many times, expensive objects should not be recreated.

Building network: creating the model using nn (2/2)

- Neural networks can be constructed using the torch.nn package.
- Forward
 - ▶ An nn.Module contains layers, and a method forward(input) that returns the output
 - ▶ You can use any of the Tensor operations in the forward function
- Backward
 - ▶ nn depends on autograd to define models and differentiate them
 - ▶ You just have to define the forward function, and the backward function (where gradients are computed) is automatically defined for you using autograd

An example of model



```
import torch
from torch import nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
        self.mp = nn.MaxPool2d(2)
        self.fc = nn.Linear(320, 10)
    def forward(self, x):
        in_size = x.size(0)
        x = F.relu(self.mp(self.conv1(x)))
        x = F.relu(self.mp(self.conv2(x)))
        x = x.view(in_size, -1) # flatten
        x = self.fc(x)
        return F.log_softmax(x)
```

Copyright (c) Ando Ki

19

Defining an optimizer

```
import torch
optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
...
for sample in dataloader:
    input = sample['input'].to(device)
    target = sample['target'].to(device)
    prediction = model(input)
    loss = loss_fn(prediction, target) # target means expected
    optimizer.zero_grad() # clears the gradients
    loss.backward()
    optimizer.step() # performs the optimization
```

Copyright (c) Ando Ki

20

Save/load models

- Saving and loading can easily be done using “torch.save” and “torch.load”
- PyTorch uses “pickling” to serialize the data.

```
>>> state = {'model_state' : model.state_dict(),
             'optimizer': optimizer.state_dict()}
>>> torch.save(state, 'state.pth')
```

```
>>> model = Model()
>>> optimizer = optim.SGD(model.parameters(), lr=0.01)
>>> checkpoint = torch.load('state.pth')
>>> model.load_state_dict(checkpoint['model_state'])
>>> optimizer.load_state_dict(checkpoint['optimizer_state'])
```

Save/load models

- First approach
 - ▶ #save only the model parameters
 - ▶ torch.save(the_model.state_dict(), PATH)
 - ▶ #load only the model parameters
 - ▶ the_model = TheModelClass(*args, **kwargs)
 - ▶ the_model.load_state_dict(torch.load(PATH))
- Second approach
 - ▶ #save the entire model
 - ▶ torch.save(the_model, PATH)
 - ▶ # load the entire model
 - ▶ the_model = torch.load(PATH)