

Install and Test Darknet on Ubuntu

2019 - 2020

Ando Ki, Ph.D.

adki@future-ds.com

Contents

- What is Darknet
- Building Darknet
- Testing Darknet using Tiny-YOLO
- Darknet profiling
- How to deal with 'gemm_nn'
- Darknet using OpenMP
- Running Tiny-YOLO with USB-CAM
- Running Tiny-YOLO with video stream

What is Darknet

- Darknet is an open source neural network framework written in C and CUDA (C ompute Unified Device Architecture) supporting CPU (Central Processing Unit) and GPU (Graphical Processing Unit) computation.

- ▶ Site: <https://pjreddie.com/darknet/>
- ▶ GitHub : <https://github.com/pjreddie/darknet>
 - ➡ This version may cause error on Raspberry Pi while running.



- Alexey's version

- ▶ <https://github.com/AlexeyAB/darknet>

"Darknet: Open Source Neural Networks in C", Joseph Redmon, <http://pjreddie.com/darknet>, 2013-2016.

3

Building Darknet

- | | |
|--|---|
| <ul style="list-style-type: none"> ■ Visit <ul style="list-style-type: none"> ▶ https://github.com/AlexeyAB/darknet ■ Download <ul style="list-style-type: none"> ▶ make a directory ▶ \$ mkdir work && cd work ▶ \$ git clone https://github.com/AlexeyAB/darknet.git ▶ \$ mv darknet darknet-alexey ■ Modify 'Makefile' <ul style="list-style-type: none"> ▶ cd darknet-alexey ▶ \$ vi Makefile <ul style="list-style-type: none"> ➡ set 1 for OpenCV if you installed it. ■ Compile <ul style="list-style-type: none"> ▶ \$ make ■ At last <ul style="list-style-type: none"> ▶ 'darknet': executable | <ul style="list-style-type: none"> ■ Nvidia CUDA related <ul style="list-style-type: none"> ▶ GPU, CUDNN, CUDNN_HALF ■ x86 Vector related <ul style="list-style-type: none"> ▶ AVX ■ Multi-core/computer related <ul style="list-style-type: none"> ▶ OpenMP ■ Shared library <ul style="list-style-type: none"> ▶ LIBSO ■ 3D camera <ul style="list-style-type: none"> ▶ ZED_CAMERA |
|--|---|

```
GPU=0
CUDNN=0
CUDNN_HALF=0
OPENCV=1
AVX=0
OPENMP=0
LIBSO=0
ZED_CAMERA=0
```

4

Darknet usage

■ \$./darknet function [function_arguments]

■ functions

- ▶ look 'darknet/examples/darknet.c' file and its related C files.
- ▶ detect [cfg_file] [weights_file] [options]
- ▶ detector [train/test/valid] [data_cfg] [cfg_file] [weights_file] [options]
- ▶ yolo [train/test/valid] [cfg_file] [weights_file] [options]
- ▶ cifar [train/test/valid] [cfg_file] [weights_file] [options]

\$./darknet detect cfg/yolov3.cfg weights/yolov3.weights data/dog.jpg

\$./darknet detector test cfg/coco.data cfg/yolov3.cfg weights/yolov3.weights data/dog.jpg

\$./darknet detector test cfg/voc.data cfg/yolo.cfg weights/yolo.weights data/dog.jpg

5

Testing Darknet using Tiny-YOLO (1/2)

■ Download weight file

- ▶ \$ cd ~/work/darknet-alexey
- ▶ \$ mkdir weights && cd weights
- ▶ for COCO trainval
 - ➔ \$ wget <https://pjreddie.com/media/files/yolov3-tiny.weights>
- ▶ for VOC2007+2012
 - ➔ wget <https://pjreddie.com/media/files/yolov2-tiny-voc.weights>

■ Run Tiny-YOLO for COCO

- ▶ \$ cd ~/work/darknet-alexey
- ▶ \$./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg

■ Run Tiny-YOLO for VOC

- ▶ \$ cd ~/work/darknet-alexey
- ▶ \$./darknet detector test cfg/voc.data cfg/yolov2-tiny-voc.cfg weights/yolov2-tiny-voc.weights data/dog.jpg


6

Testing Darknet using Tiny-YOLO (2/2)

```

pi@raspberrypi: ~/work/darknet/AlexeyAB
File Edit Tabs Help
[pi@raspberrypi] ./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg
layer  filters  size/stride  input  output
0 conv  16  3 x 3 / 1  416 x 416 x 3 -> 416 x 416 x 16 0.150 BF
1 max  2  2 x 2 / 2  416 x 416 x 16 -> 208 x 208 x 16 0.003 BF
2 conv  32  3 x 3 / 1  208 x 208 x 16 -> 208 x 208 x 32 0.399 BF
3 max  2  2 x 2 / 2  208 x 208 x 32 -> 104 x 104 x 32 0.001 BF
4 conv  64  3 x 3 / 1  104 x 104 x 32 -> 104 x 104 x 64 0.399 BF
5 max  2  2 x 2 / 2  104 x 104 x 64 -> 52 x 52 x 64 0.001 BF
6 conv  128 3 x 3 / 1  52 x 52 x 64 -> 52 x 52 x 128 0.399 BF
7 max  2  2 x 2 / 2  52 x 52 x 128 -> 26 x 26 x 128 0.000 BF
8 conv  256 3 x 3 / 1  26 x 26 x 128 -> 26 x 26 x 256 0.399 BF
9 max  2  2 x 2 / 2  26 x 26 x 256 -> 13 x 13 x 256 0.000 BF
10 conv 512 3 x 3 / 1  13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
11 max  2  2 x 2 / 1  13 x 13 x 512 -> 13 x 13 x 512 0.000 BF
12 conv 1024 3 x 3 / 1  13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
13 conv 256 1 x 1 / 1  13 x 13 x 1024 -> 13 x 13 x 256 0.089 BF
14 conv 512 3 x 3 / 1  13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
15 conv 255 1 x 1 / 1  13 x 13 x 512 -> 13 x 13 x 255 0.044 BF
16 yolo
[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
17 route 13
18 conv 128 1 x 1 / 1  13 x 13 x 256 -> 13 x 13 x 128 0.011 BF
19 upsample 2x 13 x 13 x 128 -> 26 x 26 x 128
20 route 19 8
21 conv 256 3 x 3 / 1  26 x 26 x 384 -> 26 x 26 x 256 1.196 BF
22 conv 255 1 x 1 / 1  26 x 26 x 256 -> 26 x 26 x 255 0.088 BF
23 yolo
[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
Total BFLOPS 5.571
Loading weights from weights/yolov3-tiny.weights...
Done!
data/dog.jpg: Predicted in 31128.772000 milli-seconds. 31 sec
dog: 81%
bicycle: 38%
car: 71%
truck: 41%
truck: 62%
car: 39%

```

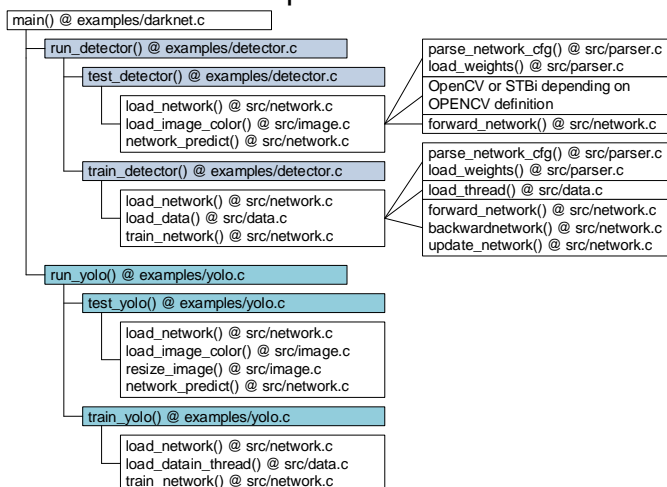


Type 'q' on the picture in order to quit.

7

Darknet example for detector

Have a look at 'darknet/examples/darknet.c' file



Darknet & YOLO

(8)

Network configuration file

■ Nodes

- ▶ [shortcut]
- ▶ [crop]
- ▶ [cost]
- ▶ [detection]
- ▶ [region]
- ▶ [local]
- ▶ **[conv] or [convolutional]**
- ▶ [deconv] or [deconvolutional]
- ▶ [activation]
- ▶ **[net] or [network]**
- ▶ [crnn]
- ▶ [gru]
- ▶ [lstm]
- ▶ [rnn]
- ▶ [conn] or [connected]
- ▶ **[max] or [maxpool]**
- ▶ [reorg]

```
[net]
batch=1
subdivisions=1
width=416
height=416
channels=3
```

```
[convolutional]
batch_normalize=1
filters=16
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

- ▶ [avg] or [avgpool]
- ▶ [dropout]
- ▶ [ln] or [normalization]
- ▶ [batchnorm]
- ▶ **[soft] or [softmax]**
- ▶ [route]

<https://github.com/cvjena/darknet/blob/master/cfg/yolo.cfg>

Darknet & YOLO

(9)

Network configuration file

■ [net] node

- ▶ batch: how many images are in each batch to average the loss over?
- ▶ subdivisions: into how many sub-batches shall each batch be divided to handle images in each sub-batch in parallel?
- ▶ height, width: input size of the network
- ▶ channels: number of components, e.g., color components
- ▶ momentum: learning parameters
- ▶ learning_rate: base learning rate
- ▶ policy: change learning rate after the corresponding steps
- ▶ steps: need to have as many steps as scale
- ▶ scales: re-scale the current learning rate by the corresponding factor once the number of steps is reached
- ▶ max_batches: max number of "iterations"
- ▶ i_snapshot_iteration: snapshot the learned weights after every k "iterations"

■ [convolutional] node

- ▶ filters: number of filters, i.e., kernels
- ▶ size: size of filter, e.g., 3 means 3x3 filter
- ▶ stride: number of stride
- ▶ pad: number of padding, e.g., 1
- ▶ activation: specify activation function

■ [maxpool] node

- ▶ size: size of filter
- ▶ stride: number of stride

■ [connected] node

- ▶ output: number of output of fully connected network
- ▶ activation: activation function

■ [detection] node or [region] node

- ▶ classes: number of classes, e.g., 20 for pascal voc (l.classes)
- ▶ coords: bounding boxes -> 4 parameters (l.n)
- ▶ side: number of cell in x and y direction
- ▶ num: number of predicted boxes per cell

Darknet & YOLO

(10)

Network configuration file

First node should be [net]

tiny-yolo.cfg

layer	filters	size	input	output
0 conv	16	3 x 3 / 1	416 x 416 x 3	416 x 416 x 16
1 max		2 x 2 / 2	416 x 416 x 16	208 x 208 x 16
2 conv	32	3 x 3 / 1	208 x 208 x 16	208 x 208 x 32
3 max		2 x 2 / 2	208 x 208 x 32	104 x 104 x 32
4 conv	64	3 x 3 / 1	104 x 104 x 32	104 x 104 x 64
5 max		2 x 2 / 2	104 x 104 x 64	52 x 52 x 64
6 conv	128	3 x 3 / 1	52 x 52 x 64	52 x 52 x 128
7 max		2 x 2 / 2	52 x 52 x 128	26 x 26 x 128
8 conv	256	3 x 3 / 1	26 x 26 x 128	26 x 26 x 256
9 max		2 x 2 / 2	26 x 26 x 256	13 x 13 x 256
10 conv	512	3 x 3 / 1	13 x 13 x 256	13 x 13 x 512
11 max		2 x 2 / 1	13 x 13 x 512	13 x 13 x 512
12 conv	1024	3 x 3 / 1	13 x 13 x 512	13 x 13 x 1024
13 conv	512	3 x 3 / 1	13 x 13 x 1024	13 x 13 x 512
14 conv	425	1 x 1 / 1	13 x 13 x 512	13 x 13 x 425
15 detection				

```
[net]
width=416
height=416
channels=3
```

```
[convolutional]
batch_normalize=1
filters=16
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=32
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=64
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=128
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=256
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=512
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=1
```

```
[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=512
activation=leaky
```

```
[convolutional]
size=1
stride=1
pad=1
filters=425
activation=linear
```

```
[region]
anchors =
```

```
0.57273,
0.677385,
1.87446, 2.06253,
3.33843, 5.47434,
7.88282, 3.52778,
9.77052, 9.16828
bias_match=1
classes=80
coords=4
num=5
softmax=1
jitter=.2
rescore=0
```

```
object_scale=5
noobject_scale=1
class_scale=1
coord_scale=1
```

```
absolute=1
thresh = .6
random=1
```

Darknet & YOLO

(11)

Network configuration file

First node should be [net]

tiny-yolo-voc.cfg

layer	filters	size	input	output
0 conv	16	3 x 3 / 1	416 x 416 x 3	416 x 416 x 16
1 max		2 x 2 / 2	416 x 416 x 16	208 x 208 x 16
2 conv	32	3 x 3 / 1	208 x 208 x 16	208 x 208 x 32
3 max		2 x 2 / 2	208 x 208 x 32	104 x 104 x 32
4 conv	64	3 x 3 / 1	104 x 104 x 32	104 x 104 x 64
5 max		2 x 2 / 2	104 x 104 x 64	52 x 52 x 64
6 conv	128	3 x 3 / 1	52 x 52 x 64	52 x 52 x 128
7 max		2 x 2 / 2	52 x 52 x 128	26 x 26 x 128
8 conv	256	3 x 3 / 1	26 x 26 x 128	26 x 26 x 256
9 max		2 x 2 / 2	26 x 26 x 256	13 x 13 x 256
10 conv	512	3 x 3 / 1	13 x 13 x 256	13 x 13 x 512
11 max		2 x 2 / 1	13 x 13 x 512	13 x 13 x 512
12 conv	1024	3 x 3 / 1	13 x 13 x 512	13 x 13 x 1024
13 conv	1024	3 x 3 / 1	13 x 13 x 1024	13 x 13 x 1024
14 conv	125	1 x 1 / 1	13 x 13 x 1024	13 x 13 x 125
15 detection				

```
[net]
subdivisions=8
width=416
height=416
channels=3
```

```
learning_rate=0.001
max_batches = 40200
policy=steps
steps=-
1,100,20000,30000
scales=.1,10,.1,.1
```

```
[convolutional]
batch_normalize=1
filters=16
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=32
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=64
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=128
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=2
```

```
[maxpool]
size=2
stride=2
```

```
[convolutional]
batch_normalize=1
filters=512
size=3
stride=1
pad=1
activation=leaky
```

```
[maxpool]
size=2
stride=1
```

```
[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
batch_normalize=1
filters=1024
size=3
stride=1
pad=1
activation=leaky
```

```
[convolutional]
size=1
stride=1
pad=1
filters=125
activation=linear
```

```
[region]
anchors =
1.08,1.19,
3.42,4.41,
6.63,11.38,
9.42,5.11,
16.62,10.52
bias_match=1
classes=20
coords=4
num=5
softmax=1
jitter=.2
rescore=1
```

```
object_scale=5
noobject_scale=1
class_scale=1
coord_scale=1
```



```
absolute=1
thresh = .6
random=1
```

Darknet & YOLO

(12)

Network visualizer

- <https://github.com/lutzroeder/netron>
- <https://lutzroeder.github.io/netron/>



13

Contents

- What is Darknet
- Building Darknet
- Testing Darknet using Tiny-YOLO
- Darknet profiling
- How to deal with 'gemm_nn'
- Darknet using OpenMP
- Running Tiny-YOLO with USB-CAM
- Running Tiny-YOLO with video stream

14

Darknet profiling (1/2)

- 'gprof' 사용
 - ▶ 컴파일 단계에서 '-pg' 선택자 사용
 - ▶ 프로그램 수행 결과로 'gmon.out' 파일 생성
 - ▶ 'gprof' 프로그램으로 분석
- 1) 'Makefile'의 'CFLAGS'에 '-pg' 추가
 - ▶ CFLAG+=-pg
- 2) 'make' 실행
 - ▶ \$ make clean && make GPROF=1
- 3) run
 - ▶ \$./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg
- 4) 'gprof' 실행
 - ▶ \$ gprof darknet gmon.out > gprof.txt
- 5) 'gprof.txt' 파일 검토
 - ▶ \$ head -20 gprof.txt

```
GPU=0
CUDNN=0
CUDNN_HALF=0
OPENCV=1
AVX=0
OPENMP=0
LIBSO=0
ZED_CAMERA=0
GPROF=0

....
ifeq ($(GPROF), 1)
CFLAGS+=-pg
endif
...
```

15

Darknet profiling (2/2)

```
File Edit Tabs Help
[pi@raspberrypi] ./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg
layer filters size/stride input output
0 conv 16 3 x 3/1 416 x 416 x 3 -> 416 x 416 x 16 0.150 BF
1 max 2 x 2/2 416 x 416 x 16 -> 208 x 208 x 16 0.003 BF
2 conv 32 3 x 3/1 208 x 208 x 16 -> 208 x 208 x 32 0.399 BF
3 max 2 x 2/2 208 x 208 x 32 -> 104 x 104 x 32 0.001 BF
4 conv 64 3 x 3/1 104 x 104 x 32 -> 104 x 104 x 64 0.399 BF
5 max 2 x 2/2 104 x 104 x 64 -> 52 x 52 x 64 0.001 BF
6 conv 128 3 x 3/1 52 x 52 x 64 -> 52 x 52 x 128 0.399 BF
7 max 2 x 2/2 52 x 52 x 128 -> 26 x 26 x 128 0.000 BF
8 conv 256 3 x 3/1 26 x 26 x 128 -> 26 x 26 x 256 0.399 BF
9 max 2 x 2/2 26 x 26 x 256 -> 13 x 13 x 256 0.000 BF
10 conv 512 3 x 3/1 13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
11 max 2 x 2/2 13 x 13 x 512 -> 13 x 13 x 512 0.000 BF
12 conv 1024 3 x 3/1 13 x 13 x 512 -> 13 x 13 x 1024 0.506 BF
13 conv 256 1 x 1/1 13 x 13 x 1024 -> 13 x 13 x 256 0.089 BF
14 conv 512 3 x 3/1 13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
15 conv 256 1 x 1/1 13 x 13 x 512 -> 13 x 13 x 256 0.044 BF
16 yolo
[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
17 route 13
18 conv 128 1 x 1/1 13 x 13 x 256 -> 13 x 13 x 128 0.011 BF
19 upsample 2x 13 x 13 x 128 -> 26 x 26 x 128
20 route 18 8
21 conv 256 3 x 3/1 26 x 26 x 384 -> 26 x 26 x 256 1.196 BF
22 conv 255 1 x 1/1 26 x 26 x 256 -> 26 x 26 x 255 0.088 BF
23 yolo
[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
Total 96.19s 5.971
Loading weights from weights/yolov3-tiny.weights...
seen 64
Done!
data/dog.jpg: Predicted in 30936.748000 milli-seconds.
dog: 01%
bicycle: 38%
car: 71%
truck: 41%
truck: 62%
car: 38%
[pi@raspberrypi] ls
dirbary/ appveyor.yml backup/ build.sh had_list
dog/ build/ build.ps1 build.sh* cfg/
cmake/ cmakeLists.txt darknet/ DarknetConfig.cmake.in
darknet_normal/ darknet_openblas* darknet_opennmp* darknet_py
darknet_video.py data/ gmm_openblas.c* gmon.out
image_yolov2.sh* image_yolov3.sh* include/ json_image_streams.sh*
LICENSE Makefile Makefile_openblas* net_cam_v3.sh*
obj/ predictions.jpg README.ad results/
scripts/ video_v2.sh* video_yolov3.sh*
weights/
[pi@raspberrypi]
```

```
File Edit Tabs Help
[pi@raspberrypi] gprof darknet gmon.out > gprof.txt
[pi@raspberrypi] head -20 gprof.txt
Flat profile:

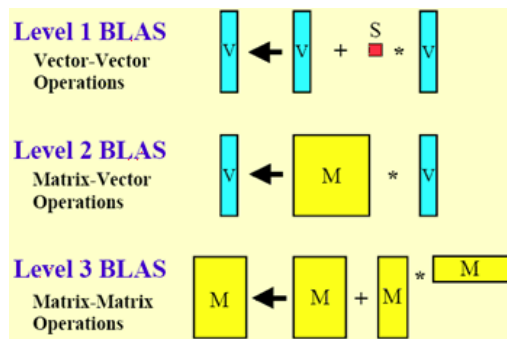
Each sample counts as 0.01 seconds.
% cumulative self self total
time seconds seconds calls s/call s/call name
96.19 30.53 30.53 3694 0.01 0.01 gemm_nn
1.20 30.91 0.38 1 0.38 0.38 fuse_conv_batchnorm
0.44 31.05 0.14 13 0.01 0.02 make_convolutional_layer
0.44 31.19 0.14 9 0.02 0.02 im2col_cpu_ext
0.43 31.33 0.14 8845488 0.00 0.00 rand_uniform
0.28 31.42 0.09 6 0.02 0.02 forward_maxpool_layer_avx
0.22 31.49 0.07 13 0.01 0.01 activate_array_cpu_custom
0.22 31.56 0.07 1 0.07 0.07 stbi_convert_format
0.13 31.60 0.04 761 0.00 0.00 mat_to_image
0.13 31.64 0.04 13 0.00 0.00 add_bias
0.13 31.68 0.04 1 0.04 0.04 resize_image
0.06 31.70 0.02 1 0.02 0.02 image_to_mat
0.06 31.72 0.02 1 0.02 0.09 save_image_options
0.03 31.73 0.01 12 0.00 0.00 activate_array
0.03 31.74 0.01 1 0.01 0.01 constrain_image
[pi@raspberrypi]
```

16

BLAS: Basic Linear Algebra Subprograms

- Level1: vector-vector operations
 - ▶ $V = V + s \times V$
- Level2: matrix-vector operations
 - ▶ $V = V + A \times V$
- Level3: matrix-matrix operations
 - ▶ $C = C + A \times B$

Level	Data Move ment	Floating-Point Ope rations	Example
Level 1	$O(N)$	$O(N)$	DDOT
Level 2	$O(N^2)$	$O(N^2)$	DGEMV
Level 3	$O(N^2)$	$O(N^3)$	DGEMM




17

GEMM of BLAS

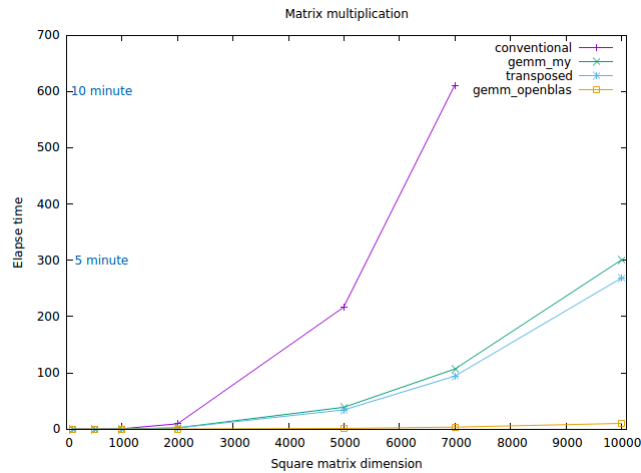
- GEMM
 - ▶ General Matrix Multiplication
 - ▶ $C = \text{BETA} \times C + \text{ALPHA} \times A \times B$
 - ▶ BLAS Level 3
 - ➔ SGEMM: single precision
 - ➔ DGEMM: double precision
 - ➔ CGEMM: single precision complex
 - ➔ ZGEMM: double precision complex

```
void gemm( int    TA
          , int    TB
          , int    M
          , int    N
          , int    K
          , float   ALPHA
          , float *A
          , int     lda
          , float   *B
          , int     ldb
          , float   BETA
          , float *C
          , int     ldc);
```



18

Square matrix multiplication



Ubuntu 16.04 / 64-bit on x86_64 / Intel Core i7-3770 CPU @ 3.4GHz x 8 / 16GigaByte


19

BLAS packages

- OpenBLAS
 - ▶ www.openblas.net
- Intel MKL (Math Kernel Library)
 - ▶ Commercial and optimized for Intel CPU
 - \$ source /opt/intel/mkl/bin/mklvars.sh intel64
- ATLAS

20

Performance comparison



21

How to deal with 'gemm_nn'(1/3)

```

#ifdef OPENBLAS
#include <cbblas.h>
void gemm( int    TA
          , int    TB
          , int    M
          , int    N
          , int    K
          , float   ALPHA
          , float  *A
          , int    lda
          , float  *B
          , int    ldb
          , float  *C
          , int    ldc)
{
    cbblas_sgemm(CblasRowMajor
                , (TA==0) ? CblasNoTrans : CblasTrans
                , (TB==0) ? CblasNoTrans : CblasTrans
                , M
                , N
                , K
                , ALPHA
                , A
                , lda
                , B
                , ldb
                , BETA
                , C
                , ldc
                , // OPENBLAS_CONST blasint ldc
                );
}

```

src/gemm.c

- Use optimized GEMM in 'src/gemm.c'
- Change 'Makefile'

```

...
OPENBLAS=0
...
ifeq ($(OPENBLAS), 1)
COMMON+=-DOPENBLAS
CFLAGS+=-DOPENBLAS -I/opt/OpenBLAS/include
LDLAGS+=-L/opt/OpenBLAS/lib/libopenblas.a
endif
...

```

Makefile

```

#else
void gemm(int TA, int TB, int M, int N, int K, float ALPHA,
          float *A, int lda,
          float *B, int ldb,
          float BETA,
          float *C, int ldc)
{
    gemm_cpu( TA, TB, M, N, K, ALPHA, A, lda, B, ldb, BETA, C, ldc);
}
#endif

```

22

How to deal with 'gemm_nn'(1/3)

```
src/gemm.c
#ifdef OPENBLAS
#include <cbblas.h>

void gemm( int    TA
, int    TB
, int    M
, int    N
, int    K
, float  ALPHA
, float  *A
, int    lda
, float  *B
, int    ldb
, float  BETA
, float  *C
, int    ldc)
{
    cbblas_sgemm(CblasRowMajor
, (TA==0) ? CblasNoTrans : CblasTrans
, (TB==0) ? CblasNoTrans : CblasTrans
, M
, N
, K
, ALPHA
, A
, lda
, B
, ldb
, BETA
, C
, ldc
, // OPENBLAS_CONST blasint M
, // OPENBLAS_CONST blasint N
, // OPENBLAS_CONST blasint K
, // OPENBLAS_CONST float alpha
, // OPENBLAS_CONST float *A
, // OPENBLAS_CONST blasint lda
, // OPENBLAS_CONST float *B
, // OPENBLAS_CONST blasint ldb
, // OPENBLAS_CONST float beta
, // float *C
, // OPENBLAS_CONST blasint ldc
);
}
```

- Use optimized GEMM in 'src/gemm.c'
- Change 'Makefile'

```
Makefile
...
OPENBLAS=1
...
ifeq ($(OPENBLAS), 1)
CFLAGS+=-DOPENBLAS
LDFLAGS+=-lopenblas
endif
...
```

```
#else
void gemm(int TA, int TB, int M, int N, int K, float ALPHA,
float *A, int lda,
float *B, int ldb,
float BETA,
float *C, int ldc)
{
    gemm_cpu( TA, TB, M, N, K, ALPHA, A, lda, B, ldb, BETA, C, ldc);
}
#endif
```

23

How to deal with 'gemm_nn'(2/3)

- Install OpenBLAS
 - ▶ \$ sudo apt-get install libopenblas-dev
- \$ make clean && make OPENBLAS=1 \
- GPROF=1
- \$./darknet detect cfg/yolov3-tiny.cfg \
- weights/yolov3-tiny.weights data/dog.jpg
- \$ gprof darknet gmon.out > gprof.txt
- \$ head -20 gprof.txt

```
pi@raspberrypi: ~/work/darknet/AlexeyAB
File Edit Tabs Help
[pi@raspberrypi] gprof darknet gmon.out > gprof.txt
[pi@raspberrypi] head -20 gprof.txt
Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self      total
time  seconds    seconds    calls   ms/call  ms/call  name
21.47    0.93     0.93         1      388.00    388.00    fuse_conv_batchnorm
13.56    1.17     0.24    8845488     0.00     0.00    rand_uniform
10.17    1.35     0.18         13     13.85    32.31    make_convolutional_layer
 3.95    1.42     0.07         13     13.85    32.31    sgemm_tcopy_L4_M4_20
 3.39    1.48     0.06         9       6.67     6.67    im2col_cpu_ext
 2.82    1.53     0.05         1      50.00    50.00    resize_image
 2.82    1.58     0.05         1      50.00    50.00    blas_thread_server
 2.26    1.62     0.04         1       0.04     0.04    inner_thread
 1.69    1.65     0.03        761     0.04     0.04    mat_to_image
 1.69    1.68     0.03        761     0.04     0.04    sgemm_kernel_L4_M4_100
 1.13    1.70     0.02        761     0.04     0.04    sgemm_kernel_L2_M4_22
 1.13    1.72     0.02        761     0.04     0.04    sgemm_kernel_L4_M4_20
 1.13    1.74     0.02        761     0.04     0.04    sgemm_ncopy_L4_M4_20
 0.56    1.75     0.01        13       0.77     0.77    add_bias
[pi@raspberrypi]
```

24

How to deal with 'gemm_nn'(3/3)

```

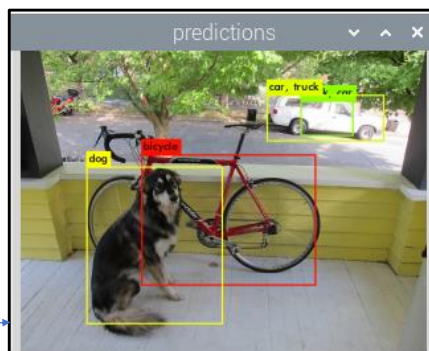
pi@raspberrypi: ~/work/darknet/AlexeyAB
File Edit Tabs Help
[pi@raspberrypi] ./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg
layer  filters  size/stride(dil)  input    output
0 conv  16          3 x 3/ 1         416 x 416 x 3 -> 416 x 416 x 16 0.150 BF
1 max   2          2 x 2/ 2         416 x 416 x 16 -> 208 x 208 x 16 0.003 BF
2 conv  32          3 x 3/ 1         208 x 208 x 16 -> 208 x 208 x 32 0.399 BF
3 max   2          2 x 2/ 2         208 x 208 x 32 -> 104 x 104 x 32 0.001 BF
4 conv  64          3 x 3/ 1         104 x 104 x 32 -> 104 x 104 x 64 0.399 BF
5 max   2          2 x 2/ 2         104 x 104 x 64 -> 52 x 52 x 64 0.001 BF
6 conv  128         3 x 3/ 1         52 x 52 x 64 -> 52 x 52 x 128 0.399 BF
7 max   2          2 x 2/ 2         52 x 52 x 128 -> 26 x 26 x 128 0.000 BF
8 conv  256         3 x 3/ 1         26 x 26 x 128 -> 26 x 26 x 256 0.399 BF
9 max   2          2 x 2/ 2         26 x 26 x 256 -> 13 x 13 x 256 0.000 BF
10 conv 512         3 x 3/ 1         13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
11 max   2          2 x 2/ 1         13 x 13 x 512 -> 13 x 13 x 512 0.000 BF
12 conv 1024        3 x 3/ 1         13 x 13 x 512 -> 13 x 13 x 1024 1.595 BF
13 conv 256         1 x 1/ 1         13 x 13 x 1024 -> 13 x 13 x 256 0.089 BF
14 conv 512         3 x 3/ 1         13 x 13 x 256 -> 13 x 13 x 512 0.399 BF
15 conv 255         1 x 1/ 1         13 x 13 x 512 -> 13 x 13 x 255 0.044 BF
16 yolo

[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
17 route 13
18 conv 128         1 x 1/ 1         13 x 13 x 256 -> 13 x 13 x 128 0.011 BF
19 upsample 2x      13 x 13 x 128 -> 26 x 26 x 128
20 route 19 8
21 conv 256         3 x 3/ 1         26 x 26 x 384 -> 26 x 26 x 256 1.196 BF
22 conv 255         1 x 1/ 1         26 x 26 x 256 -> 26 x 26 x 255 0.088 BF
23 yolo

[yolo] params: iou loss: mse, iou_norm: 0.75, cls_norm: 1.00, scale_x_y: 1.00
Total BFLOPS 5.571
Loading weights from weights/yolov3-tiny.weights...
seen 64
Done!
data/dog.jpg: Predicted in 1235.776000 milli-seconds.
dog: 81%
bicycle: 38%
car: 71%
truck: 41%
car: 39%

```

- use 'fim' to see the result
- \$ fim predict.png



25

Make a long story short

- Get Darknet-AlexeyAB version and modify (for Raspberry Pi Raspbian case)

- ▶ \$ cd ~/work/codes/darknet-projects
- ▶ \$ git clone https://github.com/AlexeyAB/darknet.git
- ▶ \$ mv darknet darknet-alexey-blas
- ▶ \$ cd darknet-alexey-blas
- ▶ \$ patch Makefile < ../patch_Makefile.txt
- ▶ \$ patch src/gemm.c < ../patch_gemm.txt
- ▶ \$ make
- ▶ \$./darknet detect cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights data/dog.jpg
- ▶ \$ fim predict.png

26

Darknet using OpenMP

- OpenMP will use multi-thread

- ▶ Install OpenMP

- \$ sudo apt-get update
 - \$ sudo apt-get install libomp-dev

```
GPU=0
CUDNN=0
CUDNN_HALF=0
OPENCV=1
AVX=0
OPENMP=1
LIBSO=0
ZED_CAMERA=0
...
```

- Simply set 'OPENMP' 1

- It can be run along with other options.

- ▶ OPENCV
 - ▶ OPENBLAS
 - ▶ GPROF

27

What is OpenCV

- OpenCV (Open Source Computer Vision Library) is an open source computer vision and machine learning software library.

- ▶ <https://opencv.org>

- What OpenCV can do :

- ▶ 1. Read and Write Images.
 - ▶ 2. Detection of faces and its features.
 - ▶ 3. Detection of shapes like Circle, rectangle etc in a image.
 - ▶ 4. Text recognition in images. (number of car license plate)
 - ▶ 5. Modifying image quality and colors
 - ▶ 6. Developing Augmented reality apps.
 - ▶ 7. Controlling camera

- Which Language it supports :

- ▶ 1. C++
 - ▶ 2. Android SDK
 - ▶ 3. Java
 - ▶ 4. Python
 - ▶ 5. C (Not recommended)



28

Installing OpenCV on Ubuntu

- If OpenCV is not installed yet, do as follows. (You need root password.)
 - ▶ \$ sudo apt-get install libopencv-dev python-opencv ffmpeg
- Set 'OPENCV' macro '1' in Makefile
 - ▶ do not forget to run 'make clean' in order to remove old files.

Makefile

```
GPU=0
CUDNN=0
OPENCV=1
OPENMP=0
DEBUG=0

# This is what I use, uncomment if you know your arch and want to specify
ARCH= -gencode arch=compute_52.code=compute_52

VPATH=./src/./examples
SLIB=libdarknet.so
ALIB=libdarknet.a
EXEC=darknet
OBJDIR=./obj/

CC=gcc
NVCC=nvcc
...
```

Set 'OPENCV' 1

29

Run YOLO with OpenCV

- Run darknet with yolo configuration with OpenCV.
 - ▶ \$./darknet detect cfg/yolov3.cfg weights/yolov3.weights data/horses.jpg


```
[adk@adki-ubuntu] ./darknet detect cfg/yolo.cfg weights/yolo.weights data/horses.jpg
layer  filters  size  input              output
0 conv  32  3 x 3 / 1  608 x 608 x 3  -> 608 x 608 x 32
1 max   2  2 x 2 / 2  608 x 608 x 32  -> 304 x 304 x 32
2 conv  64  3 x 3 / 1  304 x 304 x 32  -> 304 x 304 x 64
3 max   2  2 x 2 / 2  304 x 304 x 64  -> 152 x 152 x 64
4 conv  128 3 x 3 / 1  152 x 152 x 64  -> 152 x 152 x 128
5 conv  64  1 x 1 / 1  152 x 152 x 128 -> 152 x 152 x 64
6 conv  128 3 x 3 / 1  152 x 152 x 64  -> 152 x 152 x 128
7 max   2  2 x 2 / 2  152 x 152 x 128 -> 76 x 76 x 128
8 conv  256 3 x 3 / 1  76 x 76 x 128  -> 76 x 76 x 256
9 conv  128 1 x 1 / 1  76 x 76 x 256  -> 76 x 76 x 128
10 conv 256 3 x 3 / 1  76 x 76 x 128 -> 76 x 76 x 256
11 max   2  2 x 2 / 2  76 x 76 x 256  -> 38 x 38 x 256
12 conv 512 3 x 3 / 1  38 x 38 x 256  -> 38 x 38 x 512
13 conv 256 1 x 1 / 1  38 x 38 x 512  -> 38 x 38 x 256
14 conv 512 3 x 3 / 1  38 x 38 x 256  -> 38 x 38 x 512
15 conv 256 1 x 1 / 1  38 x 38 x 512  -> 38 x 38 x 256
16 conv 512 3 x 3 / 1  38 x 38 x 256  -> 38 x 38 x 512
17 max   2  2 x 2 / 2  38 x 38 x 512  -> 19 x 19 x 512
18 conv 1024 3 x 3 / 1  19 x 19 x 512  -> 19 x 19 x 1024
19 conv 512 1 x 1 / 1  19 x 19 x 1024 -> 19 x 19 x 512
20 conv 1024 3 x 3 / 1  19 x 19 x 512  -> 19 x 19 x 1024
21 conv 512 1 x 1 / 1  19 x 19 x 1024 -> 19 x 19 x 512
22 conv 1024 3 x 3 / 1  19 x 19 x 512  -> 19 x 19 x 1024
23 conv 1024 3 x 3 / 1  19 x 19 x 1024 -> 19 x 19 x 1024
24 conv 1024 3 x 3 / 1  19 x 19 x 1024 -> 19 x 19 x 1024
25 route 16
26 conv  64  1 x 1 / 1  38 x 38 x 512  -> 38 x 38 x 64
27 reorg  27 24  / 2  38 x 38 x 64  -> 19 x 19 x 256
28 conv 1024 3 x 3 / 1  19 x 19 x 256  -> 19 x 19 x 1024
29 conv 425 1 x 1 / 1  19 x 19 x 1024 -> 19 x 19 x 425
31 detection
mask scale: Using default '1.000000'
Loading weights from weights/yolo.weights...done!
data/horses.jpg: Predicted in 10.762035 seconds.
horse: 46%
horse: 59%
cow: 26%
horse: 91%
```

Type 'q' on the picture in order to quit.

30

Running Tiny-YOLO with USB-CAM


- `$./darknet detector demo cfg/coco.data cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights -c 0`



31

Running Tiny-YOLO with video stream

- `$./darknet detector demo cfg/coco.data cfg/yolov3-tiny.cfg weights/yolov3-tiny.weights video.mp4`



32